

Patrick Wagstrom

Software Engineering Research Scientist

patrick@wagstrom.net, <https://patrick.wagstrom.net/>

Professional Objective

A leadership position conducting research, development, and data science to generate insight and solve the world's hardest problems by incorporating multiple streams of raw data, observations of human behavior, innovative software tools, dynamic visualization, and state of the art machine learning.

Skills

Research Methods: data mining, machine learning, uncertainty analysis, qualitative interviews, social network analysis, sentiment analysis, multi-attribute utility analysis

Programming Languages: Python, JavaScript, Java, SQL, R, C

Other Assorted Skills: DevOps architecture and deployment, CloudFoundry, Kubernetes, AWS, Agile methodologies and transformation, Linux system administration

Employment History

Director of Data Science for Machine Intelligence

Capital One, New York, NY

November 2016-Present

- Architect of a platform for scalable and personalized real-time reinforcement machine learning models on top of Kubernetes, gRPC, and AWS.
- Lead a distributed team of data scientists, data engineers, and product managers to deliver our platform in a complex enterprise environment.
- Led a distributed team to build robust machine learning models that incorporate both unstructured text and structured analytical data to identify customer issues in a large corpus of customer communication.
- Manage a growing team of data scientists across four locations in three states.

Research Staff Member/Technical Lead

IBM Watson, Littleton, MA

January 2015-November 2016

- Led a globally distributed team that built the IBM Watson Conversation service - an innovative service to create rich interactions using natural language processing, entity recognition, and scripted dialog.
- Global team lead for Watson Developer Cloud Tooling. Created cutting edge applications for creating, training, and maintaining cognitive and machine learning solutions including Watson Engagement Advisor and IBM Watson Natural Language Classifier.
- Conducted research on optimal methods to build and configure cognitive and machine learning systems resulting in a 95% reduction in human time required to train these systems.
- Presented cognitive solutions and technologies to more than a dozen companies in six countries.
- Developed best-of-breed modern web application architecture based on Angular, Node.js, GitHub Enterprise, Slack, UrbanCode Deploy and more.

Research Staff Member

IBM Watson - Watson Life, Yorktown Heights, NY

January 2014-January 2015

- Leadership team member responsible for evaluating promising consumer applications of cognitive computing and guiding teams in lean startup and IBM design thinking processes to explore ideas.
- Project lead for the IBM Food Truck at SXSW which demonstrated cognitive computing to more than 4,000 people and resulted in more than 1 billion media impressions.
- Organized and directed globally distributed teams on how to contextualize solutions to cognitive computing challenges.

- Planned and executed international workshops on design thinking methodologies.
- Implemented best practices for DevOps using IBM DevOps services, IBM BlueMix, Jenkins, and more.

Research Staff Member

IBM TJ Watson Research Center, Yorktown Heights, NY

August 2009-January 2014

- Technical and strategy contributor to IBM's research and acquisition strategy for software development tools and processes.
- Analytics lead for JazzHub, IBM's cloud software development strategy. Designed analytics strategy, implemented metrics driven development, introduced A/B testing, and developed analytics dashboards.
- Developed and designed GitMiner - an open source project used by 15 universities to perform graph analysis on large scale software engineering databases such as GitHub and BitBucket.
- Led a research team to evaluate productivity of new users and small teams using IBM's enterprise software engineering and product development environments.
- Developed WhatsMyBrand, a framework for assessing an individual's personal brand by analyzing connections and contents of their actions through public social networks and relating those actions to the actions of others in their network.

Graduate Research Assistant

Carnegie Mellon University, Pittsburgh, PA

August 2003-July 2009

- Designed and developed CVSMiner an open source tool to perform social network and technical analysis of software engineering ecosystems such as GNOME and Eclipse.
- Worked with members of the GNOME Foundation and Eclipse Foundation to evaluate and improve the relationships between non-profit foundations that manage open source ecosystems and commercial firms.
- Utilized a variety of qualitative and quantitative research methods: stakeholder interviews, message analysis, natural language processing, data mining, machine learning, and social network analysis to generate insight into largely ad hoc software development processes.

Education

Ph.D. in Engineering and Public Policy and Computation, Organizations, and Society, May 2009

Carnegie Mellon University, Pittsburgh, PA

Dual degree between the Carnegie Institute of Technology (College of Engineering) and School of Computer Science at Carnegie Mellon University. Ph.D. thesis title "Vertical Interaction in Open Software Engineering Communities". Thesis advisors Dr. James Herbsleb and Dr. Kathleen Carley.

MS in Computation, Organizations, and Society, May 2007

Carnegie Mellon University, Pittsburgh, PA

Intermediate degree awarded en route to my Ph.D.

MS in Computer Science, August 2003

Illinois Institute of Technology, Chicago, IL

Focused on distributed and pervasive computing. Thesis title "Scarlet: A Framework for Context Aware Computing".

BS in Computer Science / BS in Computer Engineering / BS in Electrical Engineering, May 2002

Illinois Institute of Technology, Chicago, IL

I was on scholarship and it seemed like a good idea to keep tacking on degrees.